

CS480/680: Introduction to Machine Learning

Lec 20: Robustness

Yaoliang Yu



UNIVERSITY OF
WATERLOO

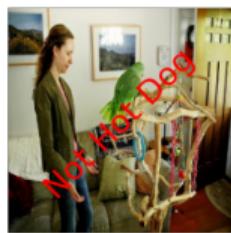
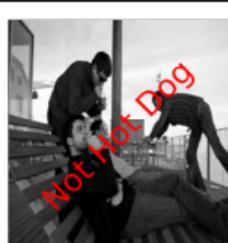
FACULTY OF MATHEMATICS
DAVID R. CHERITON SCHOOL
OF COMPUTER SCIENCE

July 17, 2024

Supervised Learning

👉 hotdog app

example results



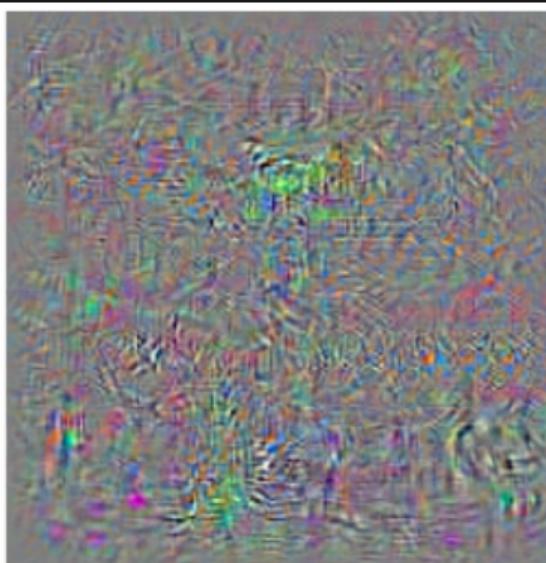
Formally

- Given a training set of pairs of examples $(\mathbf{x}_i, y_i) \in \mathbb{X} \times \mathbb{Y}$
- Return a function (classifier) $f : \mathbb{X} \rightarrow \mathbb{Y}$
- On an unseen test example \mathbf{x} , output $f(\mathbf{x})$
- The goal is to **do well** on unseen test data
 - usually do not care about performance on training data

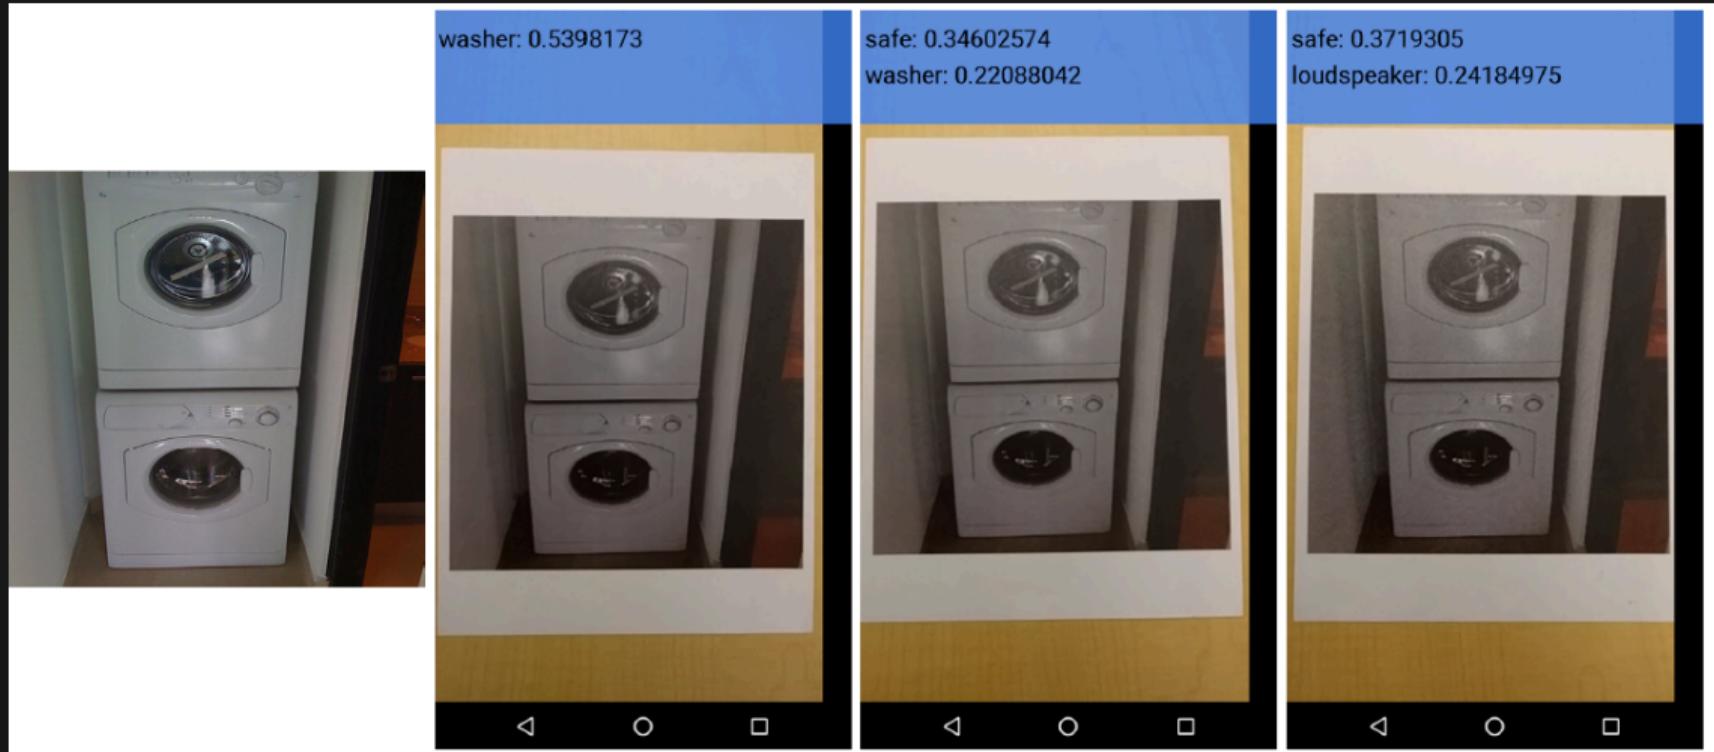
Performance Metric

- Accuracy (top-1, top-10 error, precision, recall, etc.)
- Training time
- Memory consumption
- Test (inference) time
- Robustness
- Privacy
- Fairness
- Interpretability
- Etc.

And then the Surprise







A. Kurakin, I. Goodfellow, and S. Bengio. "Adversarial examples in the physical world". In: *International Conference on Learning Representations (ICLR)*. 2017.



A. Athalye, L. Engstrom, A. Ilyas, and K. Kwok. "Synthesizing Robust Adversarial Examples". In: *ICML*. 2018.

Why Should We Care?



I. Goodfellow, P. McDaniel, and N. Papernot. "Making Machine Learning Robust Against Adversarial Inputs". *Communications of the ACM*, vol. 61, no. 7 (2018), pp. 56–66, J. Gilmer et al. "Motivating the Rules of the Game for Adversarial Example Research". 2018.



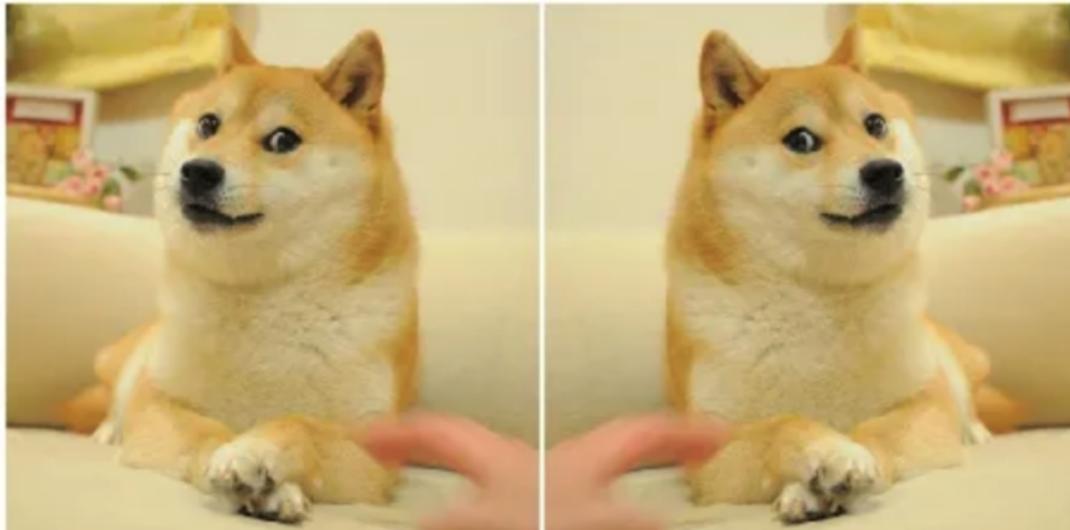
Apple's NeuralHash

- Standard hash: sensitive to small changes



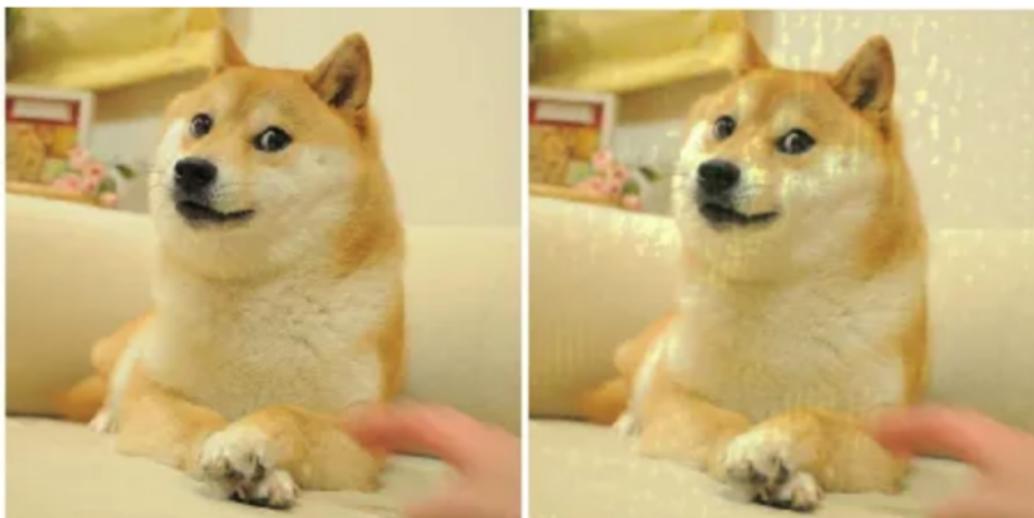
(Left) Original Doge meme, MD5: 53facff91ec83f60a88235ab628590bb | (Right) Image cropped by author,
MD5: da25273f33c4ec95f71984075079bd16

- Neural hash: contrastive training



(Left) Original Doge meme, NeuralHash: 11d9b097ac960bd2c6c131fa | (Right) Image flipped by author,
NeuralHash: 20d8f097ac960ad2c7c231fe

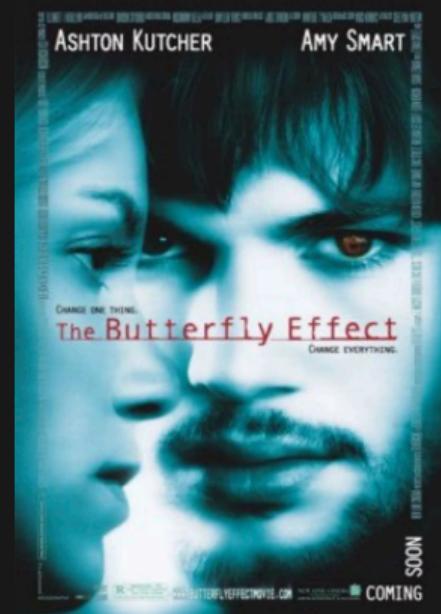
- Adversarial attack against Neural hash



(Left) Original Doge meme, NeuralHash: 11d9b097ac960bd2c6c131fa | (Right) Image generated by author,
NeuralHash: f8d1b897a45e0bf2f7e1b0fe

Formally

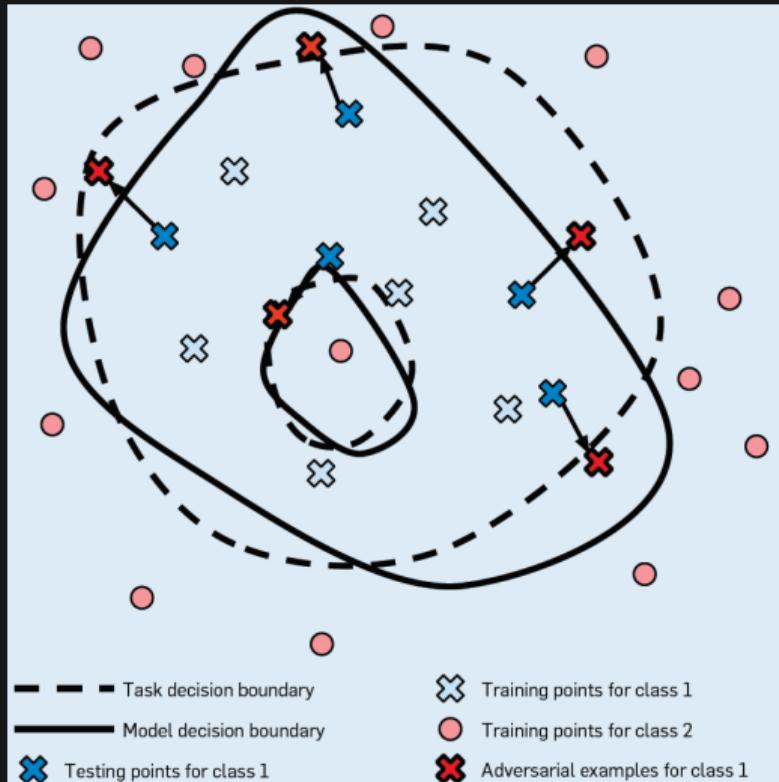
- Exist **small** Δx such that $f(x + \Delta x) \neq y(x)$
- Practically, exist **small** Δx such that $f(x + \Delta x) \neq f(x)$
 - similar if f is **very accurate**
 - such examples $x + \Delta x$ are called “**adversarial**”
- Intuitively, f is not sufficiently smooth (continuous)
- Or in fancier words, f is not **robust**
 - small perturbation should lead to small change



The Deep Challenge

$$\mathbf{x} \mapsto \varphi(\mathbf{x}) \mapsto \langle \varphi(\mathbf{x}), \mathbf{w} \rangle =: \hat{y}$$

- φ and \mathbf{w} are learned jointly



I. Goodfellow, P. McDaniel, and N. Papernot. "Making Machine Learning Robust Against Adversarial Inputs". *Communications of the ACM*, vol. 61, no. 7 (2018), pp. 56–66.

WORLD MAP



0 1,000 2,000 Miles
0 1,000 2,000 3,000 Kilometers

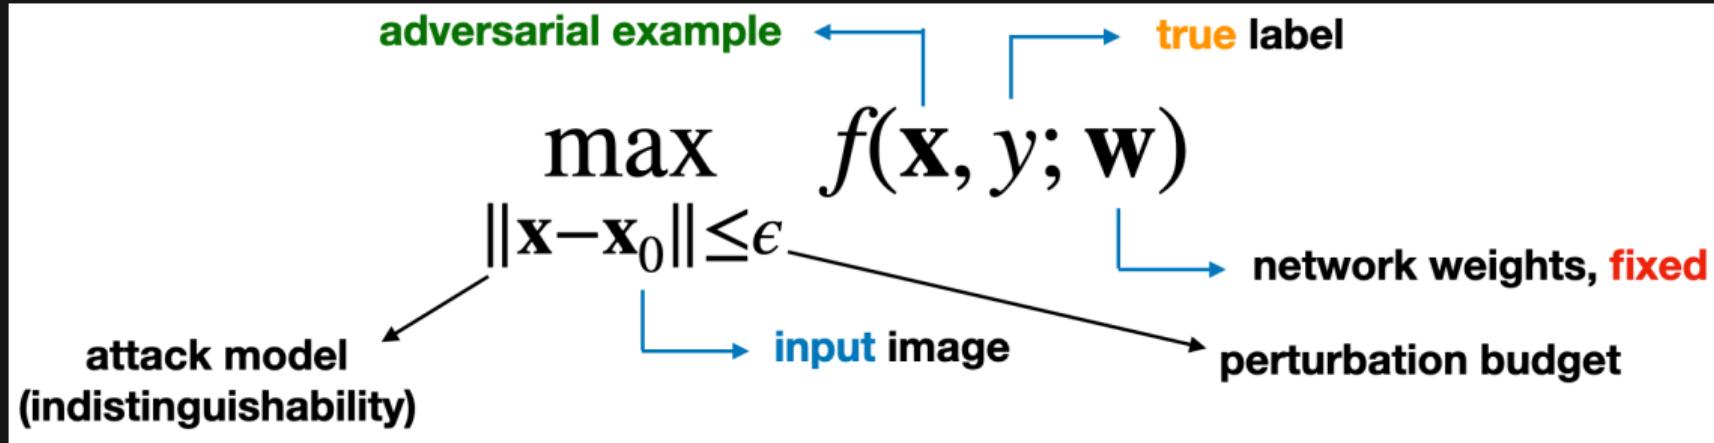
Copyright © 2019 www.mapsofworld.com

Inevitable

*“Adversarial” examples exist for **any** (non-constant) classifiers.*

- Existence is not surprising
- Ubiquity is
- How do we find adversarial examples?

Attack Formulation



- Untargeted attack, e.g. $f(\mathbf{x}, y; \mathbf{w}) = -\log p_y(\mathbf{x}; \mathbf{w})$
- Targeted attack, e.g. $f(\mathbf{x}, y; \mathbf{w}) = \log p_{\tilde{y}}(\mathbf{x}; \mathbf{w})$, for a target label $\tilde{y} \neq y$

I. J. Goodfellow, J. Shlens, and C. Szegedy. "Explaining and harnessing adversarial examples". In: *International Conference on Learning Representations (ICLR)*. 2015.

Attack Algorithms

Algorithm 1: Fast Gradient Sign Method (FGSM)

Input: model w ; data x ; label y

```
1 for  $t = 0, 1, \dots$  do
2    $x \leftarrow x + \epsilon \cdot \text{sign}(\nabla_x f(x, y; w))$            // sign grad ascent
3    $x \leftarrow P(x)$                                          // projection to stay close and valid
```

Algorithm 2: Projected Gradient Method (PGM)

Input: model w ; data x ; label y

```
1 for  $t = 0, 1, \dots$  do
2    $x \leftarrow x + \eta \cdot \nabla_x f(x, y; w)$            // grad ascent
3    $x \leftarrow P(x)$                                          // projection to stay close and valid
```

Algorithm 3: Feed-forward MLP trained with backpropogation

Input: $\mathbf{x} \in \mathbb{R}^{d_0}$, activation $\sigma : \mathbb{R} \rightarrow \mathbb{R}$, loss $\ell : \mathbb{R}^c \rightarrow \mathbb{R}$, regularizer r

```
1  $\mathbf{h}_0 \leftarrow \mathbf{x}$                                 // initialize with input data
2 for  $k = 1, \dots, l$  do                         // feature map: layer by layer
3    $\mathbf{z}_k \leftarrow W_k \mathbf{h}_{k-1} + \mathbf{b}_k$       //  $W_k \in \mathbb{R}^{d_k \times d_{k-1}}$ ,  $\mathbf{b}_k \in \mathbb{R}^{d_k}$ 
4    $\mathbf{h}_k \leftarrow \sigma(\mathbf{z}_k)$                   // element-wise
5    $\hat{\mathbf{y}} \leftarrow W_{l+1} \mathbf{h}_l + \mathbf{b}_{l+1}$     //  $W_{l+1} \in \mathbb{R}^{c \times d_l}$ ,  $\mathbf{b}_{l+1} \in \mathbb{R}^c$ 
6    $\hat{\mathbf{p}} \leftarrow \text{softmax}(\hat{\mathbf{y}})$            //  $\text{softmax}(\mathbf{a}) = \exp(\mathbf{a}) / \langle \exp(\mathbf{a}), \mathbf{1} \rangle$ 
7    $\frac{\partial \ell}{\partial W_{l+1}} \leftarrow \frac{\partial \hat{\mathbf{y}}}{\partial W_{l+1}} \cdot \frac{\partial \hat{\mathbf{p}}}{\partial \hat{\mathbf{y}}} \cdot \frac{\partial \ell}{\partial \hat{\mathbf{p}}} + \frac{\partial r}{\partial W_{l+1}}$  // we absorb  $\mathbf{b}_k$  into  $W_k$ 
8    $\frac{\partial \ell}{\partial \mathbf{h}_l} \leftarrow \frac{\partial \hat{\mathbf{y}}}{\partial \mathbf{h}_l} \cdot \frac{\partial \hat{\mathbf{p}}}{\partial \hat{\mathbf{y}}} \cdot \frac{\partial \ell}{\partial \hat{\mathbf{p}}}$  // initialize
9 for  $k = l, \dots, 1$  do                         // backward: accumulate derivatives
10   $\frac{\partial \ell}{\partial \mathbf{z}_k} \leftarrow \frac{\partial \mathbf{h}_k}{\partial \mathbf{z}_k} \cdot \frac{\partial \ell}{\partial \mathbf{h}_k}$  //  $\frac{\partial \mathbf{h}_k}{\partial \mathbf{z}_k} = \text{diag}(\sigma'(\mathbf{z}_k))$ 
11   $\frac{\partial \ell}{\partial W_k} \leftarrow \frac{\partial \mathbf{z}_k}{\partial W_k} \cdot \frac{\partial \ell}{\partial \mathbf{z}_k} + \frac{\partial r}{\partial W_k}$  //  $\frac{\partial \mathbf{z}_k}{\partial W_k} = \sum_j [\mathbf{h}_{k-1}^\top \otimes \mathbf{e}_j] \otimes \mathbf{e}_j$ 
12   $\frac{\partial \ell}{\partial \mathbf{h}_{k-1}} \leftarrow \frac{\partial \mathbf{z}_k}{\partial \mathbf{h}_{k-1}} \cdot \frac{\partial \ell}{\partial \mathbf{z}_k}$  //  $\frac{\partial \mathbf{z}_k}{\partial \mathbf{h}_{k-1}} = W_k^\top$ 
```

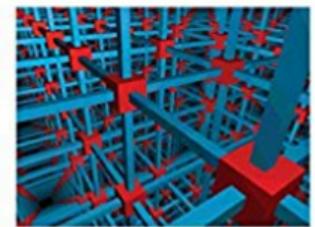
Adversarial Training

$$\min_{\mathbf{w}} \hat{\mathbb{E}} \left[\underbrace{\max_{\|\Delta \mathbf{x}\| \leq \epsilon} \ell(\mathbf{x} + \Delta \mathbf{x}; \mathbf{w})}_{\ell_\epsilon(\mathbf{x}; \mathbf{w})} \right]$$

- $\epsilon = 0$ reduces to the usual training
- $\epsilon \neq 0$ amounts to a change in the loss
- A min-max game between the defender (model \mathbf{w}) and attacker (adversarial example $\mathbf{x} + \Delta \mathbf{x}$)
- SGD on \mathbf{w} where the inner max is solved by an attack algorithm

Princeton Series in APPLIED MATHEMATICS

Robust Optimization



Aharon Ben-Tal
Laurent El Ghaoui
Arkadi Nemirovski

Lasso Revisited

- Consider the linear regression problem:

$$\min_{\mathbf{w} \in \mathbb{R}^d} \|X\mathbf{w} - \mathbf{y}\|_2, \quad \text{where } X \in \mathbb{R}^{n \times d}, \mathbf{y} \in \mathbb{R}^n$$

- Suppose we perturb each *feature* independently, and we are interested in solving the robust linear regression problem:

$$\min_{\mathbf{w} \in \mathbb{R}^d} \max_{\forall j, \|\mathbf{z}_j\|_2 \leq \lambda} \|(X + Z)\mathbf{w} - \mathbf{y}\|_2, \quad \text{where } Z = [\mathbf{z}_1, \dots, \mathbf{z}_d] \in \mathbb{R}^{n \times d}$$

- Prove that robust linear regression is exactly equivalent to (square-root) Lasso:

$$\min_{\mathbf{w} \in \mathbb{R}^d} \|X\mathbf{w} - \mathbf{y}\|_2 + \lambda \|\mathbf{w}\|_1, \quad \text{where } \|\mathbf{w}\|_1 = \sum_j |w_j|$$

Robustness More Broadly

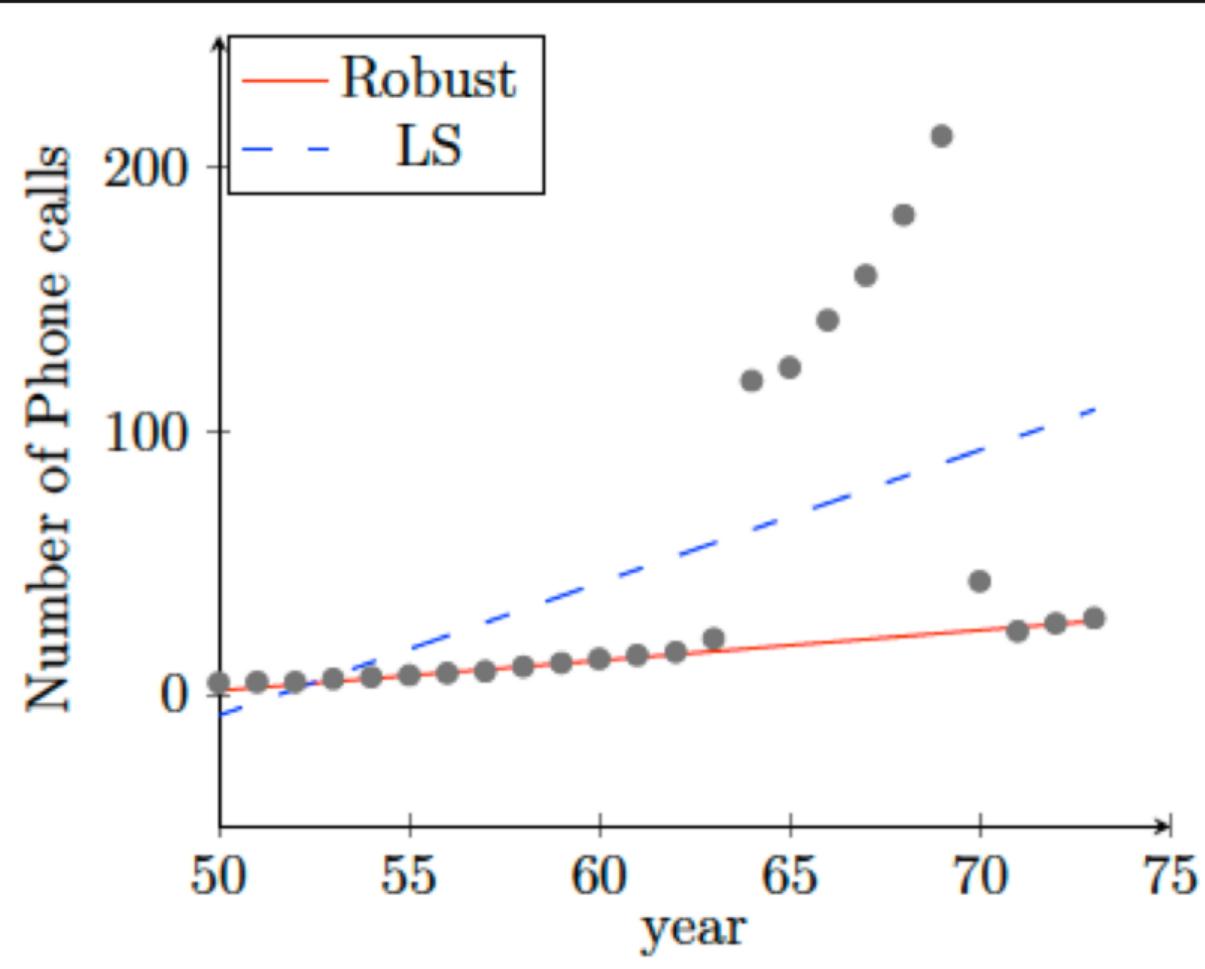
$$\|X\mathbf{w} - \mathbf{y}\|_2^2 \quad vs. \quad \|X\mathbf{w} - \mathbf{y}\|_1$$

- Least-squares (Gauss) vs. least absolute deviation (Laplace)

S. PORTNOY AND R. KOENKER



S. Portnoy and R. Koenker. "The Gaussian hare and the Laplacian tortoise: computability of squared-error versus absolute-error estimators". *Statistical Science*, vol. 12, no. 4 (1997), pp. 279–300.



WILEY SERIES IN PROBABILITY AND STATISTICS

Robust Statistics

Second Edition



Peter J. Huber
Elvezio M. Ronchetti

WILEY

WILEY

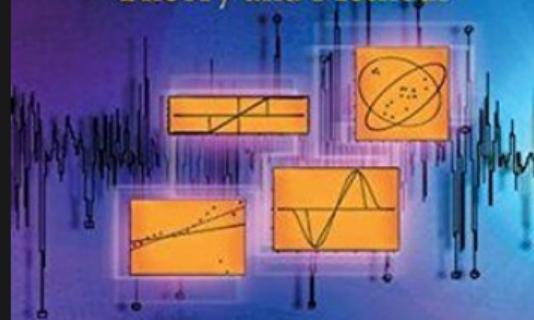
Robust Statistics The Approach Based on Influence Functions

FRANK R. HAMPEL
ELVEZIO M. RONCHETTI
PETER J. ROUSSEEUW
WERNER A. STAHEL

WILEY SERIES IN PROBABILITY AND STATISTICS

Robust Statistics

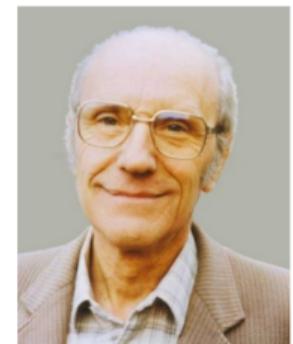
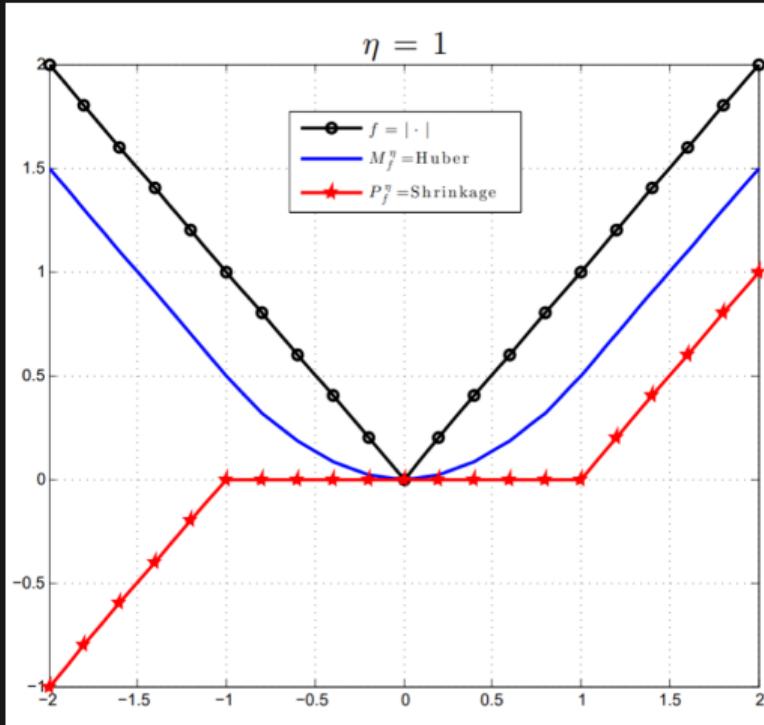
Theory and Methods



Ricardo Maronna, Doug Martin
and Victor Yohai

WILEY SERIES IN PROBABILITY AND STATISTICS

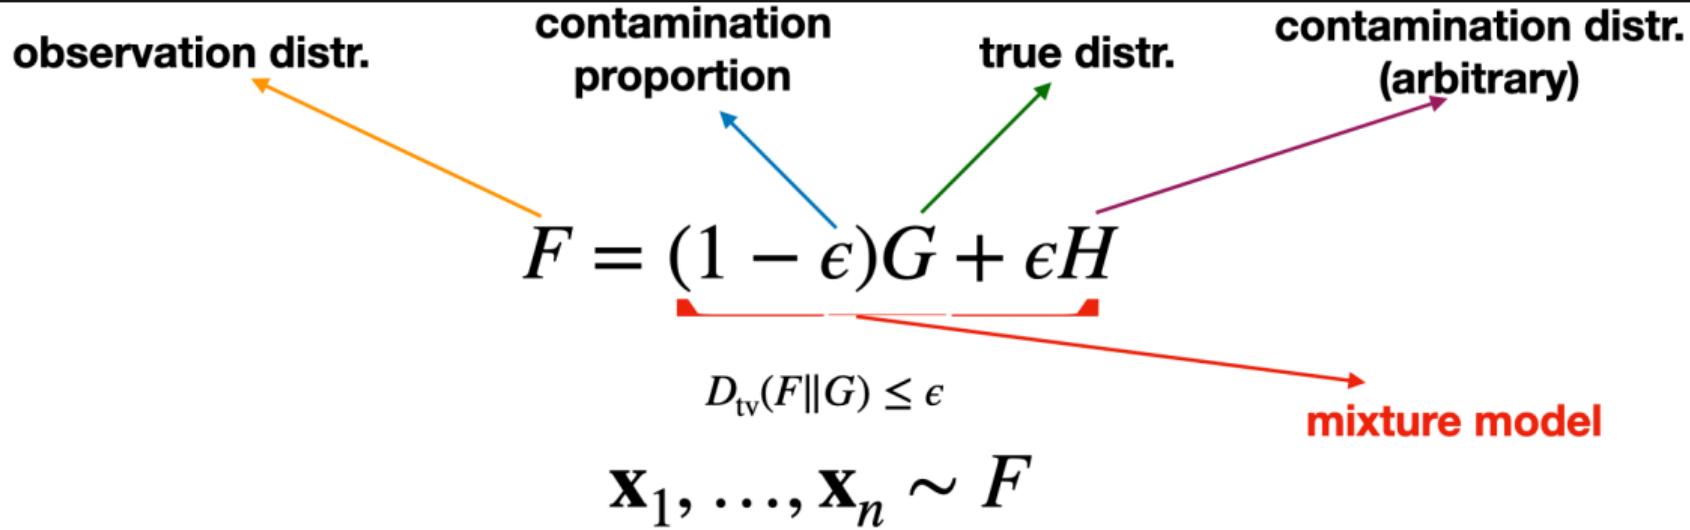
Huber's loss is Moreau's envelope



$$\begin{aligned} M_{|\cdot|}^\eta(t) &= \min_s \frac{1}{2}(s-t)^2 + \eta|s| \\ &= \begin{cases} \frac{1}{2}t^2, & |t| \leq \eta \\ \eta|t| - \frac{1}{2}\eta^2, & |t| \geq \eta \end{cases} \end{aligned}$$

J. J. Moreau. "Fonctions convexes duales et points proximaux dans un espace hilbertien". *Comptes rendus hebdomadaires des séances de l'Académie des sciences*, vol. 255 (1962), pp. 2897–2899, P. J. Huber. "Robust estimation of a location parameter". *Annals of Statistics*, vol. 35, no. 1 (1964), pp. 73–101.

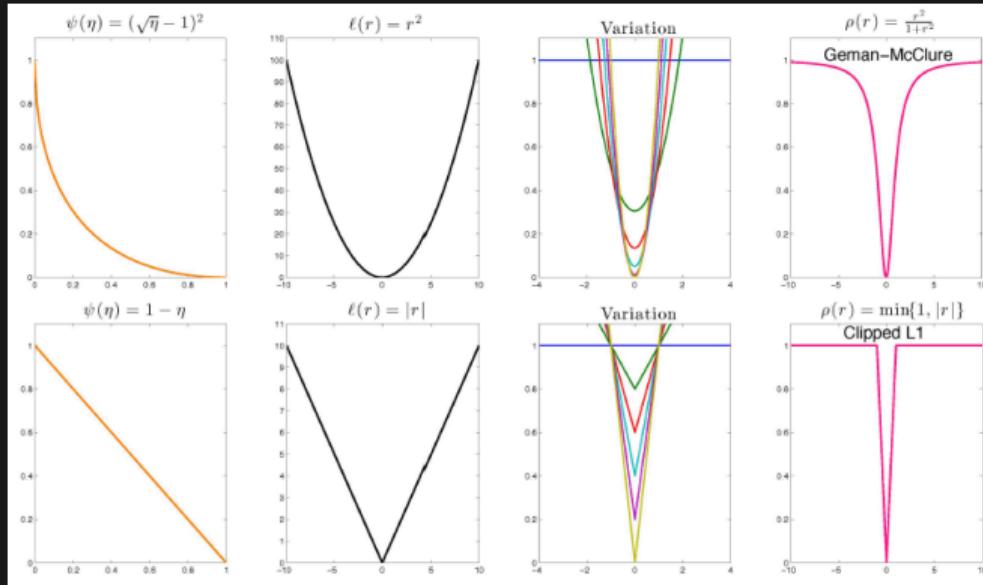
Huber's ϵ -contamination Model



- With probability (w.p.) ϵ , \mathbf{x}_i is from contamination distribution H
- W.p. $1 - \epsilon$, \mathbf{x}_i from true distribution G
- Roughly ϵ proportion is (arbitrarily) contaminated
- A mixture of two distributions: don't know which data comes from which distr.

Variational Loss

$$\tilde{\ell}(t) := \min_{\eta \in [0,1]} \eta \cdot \ell(t) + \psi(\eta)$$



M. J. Black and A. Rangarajan. "On the unification of line processes, outlier rejection, and robust statistics with applications in early vision". *International Journal of Computer Vision*, vol. 19 (1996), pp. 57–91, Y. Yu, Ö. Aslan, and D. Schuurmans. "A Polynomial-time Form of Robust Regression". In: *Advances in Neural Information Processing Systems 26 (NIPS)*. 2012.

