

CS480/680: Introduction to Machine Learning
Homework 3

Due: 11:59 pm, July 08, 2024, submit on LEARN.


NAME
student number

Submit your writeup in pdf and all source code in a zip file (with proper documentation). Write a script for each programming exercise so that the TA can easily run and verify your results. Make sure your code runs!
[Text in square brackets are hints that can be ignored.]


Exercise 1: Vision Transformers (10 pts)

Please follow the instructions of this [ipynb file](#).

- (1+3+2 = 6 pts) Complete the missing coding parts in the provided [ipynb file](#).
- (1 pt) Visualization of patches:



- (1 pt) The test accuracy I obtained on MNIST is: xxx%
- (2 pts) Training / Validation accuracy vs. epoch:



Exercise 2: Adaboost (8 pts)

In this exercise we will implement Adaboost. Recall that Adaboost aims at minimizing the exponential loss:

$$\min_{\mathbf{w}} \sum_i \exp \left(-y_i \sum_j w_j h_j(\mathbf{x}_i) \right), \quad (1)$$

where h_j are the so-called weak learners, and the combined classifier

$$h_{\mathbf{w}}(\mathbf{x}) := \sum_j w_j h_j(\mathbf{x}). \quad (2)$$

Note that we **assume** $y_i \in \{\pm 1\}$ in this exercise, and we **simply take** $h_j(\mathbf{x}) = \text{sign}(\pm x_j + b_j)$ for some $b_j \in \mathbb{R}$. Upon **defining** $M_{ij} = y_i h_j(\mathbf{x}_i)$, we may simplify our problem further as:

$$\min_{\mathbf{w} \in \mathbb{R}^d} \mathbf{1}^\top \exp(-M\mathbf{w}), \quad (3)$$

where \exp is applied component-wise and $\mathbf{1}$ is the vector of all 1s.

Recall that $(s)_+ = \max\{s, 0\}$ is the positive part while $(s)_- = \max\{-s, 0\} = |s| - s_+$.

Algorithm 1: Adaboost.

Input: $M \in \mathbb{R}^{n \times d}$, $\mathbf{w}_0 = \mathbf{0}_d$, $\mathbf{p}_0 = \mathbf{1}_n$, $\text{max_pass} = 300$

Output: \mathbf{w}

```

1 for  $t = 0, 1, 2, \dots, \text{max\_pass}$  do
2    $\mathbf{p}_t \leftarrow \mathbf{p}_t / (\mathbf{1}^\top \mathbf{p}_t)$  // normalize
3    $\boldsymbol{\epsilon}_t \leftarrow (M)^\top_- \mathbf{p}_t$  //  $(\cdot)_-$  applied component-wise
4    $\boldsymbol{\gamma}_t \leftarrow (M)^\top_+ \mathbf{p}_t$  //  $(\cdot)_+$  applied component-wise
5    $\boldsymbol{\beta}_t \leftarrow \frac{1}{2}(\ln \boldsymbol{\gamma}_t - \ln \boldsymbol{\epsilon}_t)$  //  $\ln$  applied component-wise
6   choose  $\boldsymbol{\alpha}_t \in \mathbb{R}^d$  // decided later
7    $\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + \boldsymbol{\alpha}_t \odot \boldsymbol{\beta}_t$  //  $\odot$  component-wise multiplication
8    $\mathbf{p}_{t+1} \leftarrow \mathbf{p}_t \odot \exp(-M(\boldsymbol{\alpha}_t \odot \boldsymbol{\beta}_t))$  //  $\exp$  applied component-wise
```

1. (2 pts) We claim that Algorithm 1 is indeed the celebrated Adaboost algorithm if the following holds:

- $\boldsymbol{\alpha}_t$ is one-hot (i.e., 1 at some entry and 0 everywhere else), namely, it indicates which weak classifier is chosen at iteration t .
- $M \in \{\pm 1\}^{n \times d}$, i.e., if all weak classifiers are $\{\pm 1\}$ -valued.

With the above conditions, prove that (a) $\boldsymbol{\gamma}_t = \mathbf{1} - \boldsymbol{\epsilon}_t$, and (b) the equivalence between Algorithm 1 and the Adaboost algorithm in class. [Note that our labels here are $\{\pm 1\}$ and our \mathbf{w} may have nothing to do with the one in class.]

Ans:

2. (2 pts) Let us derive each week learner h_j . Consider each feature in turn, we train d linear classifiers that each aims to minimize the weighted training error:

$$\min_{b_j \in \mathbb{R}, s_j \in \{\pm 1\}} \sum_{i=1}^n p_i \mathbb{I}[y_i(s_j x_{ij} + b_j) \leq 0], \quad (4)$$

where the weights $p_i \geq 0$ and $\sum_i p_i = 1$. Find (with justification) an optimal value for each b_j and s_j . [If multiple solutions exist, you can use the middle value.] If it helps, you may assume p_i is uniform, i.e., $p_i \equiv \frac{1}{n}$.

Ans:

3. (2 pts) [Parallel Adaboost.] Implement Algorithm 1 with the following choices:
- $\alpha_t \equiv 1$
 - pre-process M by dividing a constant so that for all i (row), $\sum_j |M_{ij}| \leq 1$.


Run your implementation on the **default** dataset (available on **course website**), and report the training loss in (3), training error, and test error w.r.t. the iteration t , where

$$\text{error}(\mathbf{w}; \mathcal{D}) := \frac{1}{|\mathcal{D}|} \sum_{(\mathbf{x}, y) \in \mathcal{D}} \mathbb{I}[yh_{\mathbf{w}}(\mathbf{x}) \leq 0]. \quad (5)$$

[Recall that $h_{\mathbf{w}}(\mathbf{x})$ is defined in (2) while each h_j is decided in Ex 2.2 with uniform weight $p_i \equiv \frac{1}{n}$. In case you fail to determine h_j , in Ex 2.3 and Ex 2.4 you may simply use $h_j(\mathbf{x}) = \text{sign}(x_j - m_j)$ where m_j is the median value of the j -th feature in the training set.]

[Note that \mathbf{w}_t is dense (i.e., using all weak classifiers) even after a single iteration.]

Ans: We report all 3 curves in one figure, with clear coloring and legend to indicate which curve is which.



4. (2 pts) [Sequential Adaboost.] Implement Algorithm 1 with the following choice:

- $j_t = \text{argmax}_j |\sqrt{\epsilon_{t,j}} - \sqrt{\gamma_{t,j}}|$ and α_t has 1 on the j_t -th entry and 0 everywhere else.
- pre-process M by dividing a constant so that for all i and j , $|M_{ij}| \leq 1$.

Run your implementation on the **default** dataset (available on **course website**), and report the training loss in (3), training error, and test error in (5) w.r.t. the iteration t .

[Note that \mathbf{w}_t has at most t nonzeros (i.e., weak classifiers) after t iterations.]

Ans: We report all 3 curves in one figure, with clear coloring and legend to indicate which curve is which.

